

Deep Learning HomeWork 02

Mahsa Naseri 402209015

بخش تئوری

سوال ۱:

$$1 \quad E_{w,b} = - \sum_n y_n \ln \hat{y}(x_n) + (1-y_n) \ln (1-\hat{y}(x_n))$$

Sigmoid \rightarrow Activate Function $f(x) = \sigma(x) = \frac{1}{1+e^{-x}}$

$$f'(x) = f(x)(1-f(x))$$

$$\hat{y}(x_n) = \sigma(w \cdot x_n + b) = \frac{1}{1+e^{-(w \cdot x_n + b)}}$$

$$\star \frac{\partial E_{w,b}}{\partial w_j} = - \sum_n \left(\frac{\partial}{\partial w_j} (y_n \ln(\hat{y}(x_n))) + \frac{\partial}{\partial w_j} ((1-y_n) \ln(1-\hat{y}(x_n))) \right)$$

$$1 \quad \frac{\partial}{\partial w_j} (y_n \ln \hat{y}(x_n)) = y_n \frac{1}{\hat{y}(x_n)} \cdot \frac{\partial \hat{y}(x_n)}{\partial w_j} = \frac{y_n}{\hat{y}(x_n)} (\hat{y}(x_n)(1-\hat{y}(x_n))) x_{nj}$$

$$= y_n (1-\hat{y}(x_n)) x_{nj}$$

$$2 \quad \frac{\partial}{\partial w_j} ((1-y_n) \ln(1-\hat{y}(x_n))) = (1-y_n) \left(\frac{1}{1-\hat{y}(x_n)} \right) \cdot \frac{\partial (1-\hat{y}(x_n))}{\partial w_j}$$

$$= \frac{1-y_n}{1-\hat{y}(x_n)} \cdot \frac{- (\hat{y}(x_n)(1-\hat{y}(x_n))) x_{nj}}{1} = -\hat{y}(x_n)(1-y_n)(-x_{nj})$$

$$\frac{\partial E_{w,b}}{\partial w_j} = - \sum_n (y_n (1-\hat{y}(x_n)) x_{nj}) + (\hat{y}(x_n)(1-y_n)(-x_{nj}))$$

$$= - \sum_n x_{nj} (-y_n (1-\hat{y}(x_n)) + \hat{y}(x_n)(1-y_n))$$

$$= \sum_n x_{nj} (\hat{y}(x_n) - y_n)$$

$$\frac{\partial E_{w,b}}{\partial w_j} = 0 \Rightarrow \sum_n x_{nj} (\hat{y}(x_n) - y_n) = 0 \xrightarrow{x_{nj}=0 \quad x} \hat{y}(x_n) = y_n \quad \checkmark$$

$$\sum_n \left(\frac{1}{1+e^{-(w \cdot x_n + b)}} - y_n \right) x_n = 0$$

Gradient Descent: $w_t = w_{t-1} - \eta \frac{\partial E}{\partial w}$

$$w_t = w_{t-1} - \eta \sum_n (\hat{y}(x_n) - y_n) x_n$$

$$\frac{\partial E_{w,b}}{\partial b} = - \sum \left(\frac{\partial}{\partial b} (y_n \ln \hat{y}(x_n)) + \frac{\partial}{\partial b} ((1-y_n) \ln (1-\hat{y}(x_n))) \right)$$

$$1: \frac{\partial}{\partial b} (y_n \ln \hat{y}(x_n)) = y_n \frac{1}{\hat{y}(x_n)} \frac{\partial \hat{y}(x_n)}{\partial b} = \frac{y_n}{\hat{y}(x_n)} (\hat{y}(x_n)(1-\hat{y}(x_n)))$$

$$= y_n (1-\hat{y}(x_n))$$

$$2: \frac{\partial}{\partial b} ((1-y_n) \ln (1-\hat{y}(x_n))) = \frac{1-y_n}{1-\hat{y}(x_n)} \cdot \frac{\partial (1-\hat{y}(x_n))}{\partial b} = -\frac{(1-y_n)}{1-\hat{y}(x_n)} (\hat{y}(x_n)(1-\hat{y}(x_n)))$$

$$\frac{\partial E_{w,b}}{\partial b} = - \sum (y_n (1-\hat{y}(x_n)) + \hat{y}(x_n) (y_n - 1)) = (y_n - 1) \hat{y}(x_n)$$

$$= - \sum_n (y_n - \hat{y}(x_n)) = 0 \quad [y_n = \hat{y}(x_n)]$$

$$\frac{\partial E_{w,b}}{\partial b} = 0 \xrightarrow{GD} b_t = b_{t-1} - \eta \frac{\partial E}{\partial b} = b_{t-1} + \eta \sum_n (y_n - \hat{y}(x_n))$$

$$b_t = b_{t-1} + \eta \sum_n (y_n - \hat{y}(x_n))$$

min تابع هزینه زمانی را می‌خواهیم به حداقل برسانیم

2

convariate shift: این پدیده زمانی رخ می‌دهد که توزیع داده‌های ورودی در لایه‌های مختلف شبکه تغییر می‌کند. وقتی پارامترهای لایه‌ی قبل در طول آموزش تغییر می‌کند، در واقع توزیع داده‌های ورودی لایه‌های بعدی نیز تغییر می‌کند. به این پدیده را convariate shift می‌گویند. به اجبار خودش را تصحیح دهد و این باعث می‌شود آموزش و بهتری می‌شود.

BN: به نرمال سازی ورودی‌ها (میانگین و واریانس) می‌گویند و این کار را برای هر Batch (ایتم‌های هر) باید می‌شود توزیع داده‌ها باید نرمال شده شود. از تغییرات پدید در توزیع اینها جلوگیری می‌کند. به این افزایش سرعت یادگیری را می‌توان گفت.

از learning rate بالاتر می‌شود و همچنین از مشکل gradient vanishing جلوگیری می‌کند.

$$\mu = \frac{1}{n} \sum_{j=1}^n x_j, \quad \sigma^2 = \frac{1}{n} \sum_{j=1}^n (x_j - \mu)^2$$

می‌تواند x batch of input

normalize $\rightarrow \hat{x}_i = \frac{x_i - \mu}{\sqrt{\sigma^2 + \epsilon}}$ برای پایداری عددی

(ب) هر batch عملیات‌های تصادفی از داده‌ها است، همچنین محاسبه می‌تواند داریانی برای هر batch نیز توزیع کوچک به سبب اضافه می‌کند و با اضافه کردن توزیع هر Batch باعث نرمال‌تر کردن مانند یک regularizer عمل کرده و از overfitting جلوگیری می‌کند و به دلیل اجازه می‌دهد به یادگیری ویژگی‌های ثابت در mini-batch‌های مختلف تقسیم‌بندی تجربه داشته باشد. پس جلوگیری از overfitting که و همچنین مدل را روی داده‌های دیده نشده بهبود می‌بخشد.

$$\begin{cases} \hat{x}_i = x_i - \mu & [\mu_1, \mu_2, \dots, \mu_n] & \mu = \frac{1}{n} \sum_{j=1}^n x_j & n \rightarrow \text{mini-batch} \\ y_i = \gamma \hat{x}_i + \beta & [y_1, y_2, \dots, y_n] & \frac{\partial L}{\partial x_i} = ? & \text{Cost } F = L \end{cases}$$

chain rule: $\frac{\partial L}{\partial \hat{x}_i} = \sum_{j=1}^n \frac{\partial L}{\partial y_j} \cdot \frac{\partial y_j}{\partial \hat{x}_i} \stackrel{\textcircled{1}}{=} \frac{\partial L}{\partial \hat{x}_i} = \gamma \cdot \frac{\partial L}{\partial y_i} \quad \textcircled{\text{I}}$

$$\left. \begin{aligned} \frac{\partial y_j}{\partial \hat{x}_i} &= \gamma & \frac{\partial y_j}{\partial \hat{x}_i} &= \gamma \text{ in the same index } [j=i] \\ & & \frac{\partial y_j}{\partial \hat{x}_i} &= 0 \quad j \neq i \end{aligned} \right\} \textcircled{1}$$

$$\begin{aligned} \hat{x}_i &= x_i - \mu, \quad \mu = \frac{1}{n} \sum_{j=1}^n x_j \\ \left[\begin{aligned} \frac{\partial \hat{x}_i}{\partial x_i} &= 1 - \frac{1}{n} & [j=i] \\ \frac{\partial \hat{x}_i}{\partial x_i} &= -\frac{1}{n} & [j \neq i] \end{aligned} \right] \quad \textcircled{\text{II}} \end{aligned}$$

$$\textcircled{\text{I}}, \textcircled{\text{II}} : \frac{\partial L}{\partial x_i} = \sum_{j=1}^n \frac{\partial L}{\partial \hat{x}_j} \cdot \frac{\partial \hat{x}_j}{\partial x_i}$$

$$\frac{\partial L}{\partial x_i} = \frac{\partial L}{\partial \hat{x}_i} \left(1 - \frac{1}{n}\right) + \sum_{j \neq i} \frac{\partial L}{\partial \hat{x}_j} \left(-\frac{1}{n}\right)$$

$$\boxed{\frac{\partial L}{\partial x_i} = \gamma \frac{\partial L}{\partial y_i} \left(1 - \frac{1}{n}\right) - \frac{\gamma}{n} \sum \frac{\partial L}{\partial y_i}}$$

ت) $n=1, n \rightarrow \infty$

$$n \rightarrow \infty \begin{cases} 1 - \frac{1}{n} \approx 1 \\ -\frac{1}{n} \approx 0 \end{cases} \Rightarrow \frac{\partial L}{\partial x_i} \approx \gamma \cdot \frac{\partial L}{\partial y_i}$$

$$n=1 \begin{cases} 1 - \frac{1}{n} = 0 \\ -\frac{1}{n} = -1 \end{cases} \rightarrow \frac{\partial L}{\partial x_i} = \gamma \cdot \frac{\partial L}{\partial y_i}$$

برای $n=1$ $\mu = x_1 = 0$ $\hat{x}_1 = x_1 - \mu = 0$
 آیا می‌توانیم برای $n > 1$ به جای μ از x_1 استفاده کنیم؟
 در واقع؟ به ازای $n=1$ $\mu = x_1$ و برای $n > 1$ $\mu \neq x_1$ است.

اما برای $n=1$ $\hat{x}_i = 0$ \rightarrow نیازی به محاسبه \hat{x}_i نیست.

3

$$\frac{\partial \hat{y}_k}{\partial z_i^{(2)}} = ?$$

$$z^{(1)} = w^{(1)} x + b^{(1)}, \quad a^{(1)} = \text{leakyRelu}(z^{(1)})$$

$$z^{(r)} = w^{(2)} a^{(1)} + b^{(r)}, \quad \hat{y} = \text{softmax}(z^{(2)})$$

$$\text{softmax: } \hat{y}_k = \frac{e^{z_k^{(r)}}}{\sum_j e^{z_j^{(r)}}}$$

$$\frac{\partial \hat{y}_k}{\partial z_i^{(r)}} = \begin{cases} i=k & \frac{\partial \hat{y}_k}{\partial z_i^{(r)}} = \hat{y}_k (1 - \hat{y}_k) \\ i \neq k & \frac{\partial \hat{y}_k}{\partial z_i^{(r)}} = \frac{-e^{z_k^{(r)}} e^{z_i^{(r)}}}{(\sum_{j=1}^K e^{z_j^{(r)}})^2} = -\hat{y}_k \hat{y}_i \end{cases}$$

صاف و صاف
↑
 $\frac{\partial \hat{y}_k}{\partial z_i^{(r)}} = \hat{y}_k (\delta_{ki} - \hat{y}_i)$

$$L = \sum_{i=1}^K -y_i \log \hat{y}_i$$

$$\frac{\partial L}{\partial z_i^{(r)}} = ?$$

$$y_k = 1 \quad y_j = 0 \quad (i \neq j)$$

$$\frac{\partial L}{\partial z_i^{(2)}} = \frac{\partial L}{\partial \hat{y}_i} \cdot \frac{\partial \hat{y}_i}{\partial z_i^{(2)}}$$

$$\frac{\partial L}{\partial z_i^{(r)}} = \frac{\partial}{\partial \hat{y}_i} \left(\sum_{i=1}^K -y_i \log \hat{y}_i \right) \cdot (-\hat{y}_k \hat{y}_i)$$

$-y_k \log \hat{y}_k = -\log \hat{y}_k$

$$\frac{\partial L}{\partial z_i^{(r)}} = \frac{\partial (-\log \hat{y}_k)}{\partial \hat{y}_i} \cdot (-\hat{y}_k \hat{y}_i) = \frac{-1}{\hat{y}_k} \cdot (-\hat{y}_k \hat{y}_i) = \hat{y}_i$$

$$\frac{\partial L}{\partial w^{(1)}} = \frac{\partial L}{\partial z^{(r)}} \cdot \frac{\partial z^{(r)}}{\partial a^{(1)}} \cdot \frac{\partial a^{(1)}}{\partial w^{(1)}} \quad \textcircled{I} \quad \text{بـ}$$

$$\frac{\partial L}{\partial z^{(r)}} = \hat{y} \quad \text{لنزيارت صلي}$$

$$\frac{\partial z^{(r)}}{\partial a^{(1)}} = w^{(2)} \quad \star \quad \frac{\partial L}{\partial a^{(1)}} = \hat{y} w^{(2)} \quad \textcircled{II}$$

$$\hat{a}_j^{(1)} = \begin{cases} z_j^{(1)} & z_j^{(1)} > 0 \\ 0 & z_j^{(1)} < 0 \end{cases} \Rightarrow \frac{\partial \hat{a}^{(1)}}{\partial z^{(1)}} = \begin{cases} 1 & z^{(1)} > 0 \\ 0 & z^{(1)} < 0 \end{cases}$$

$$\frac{\partial a^{(1)}}{\partial w^{(1)}} = \frac{\partial a^{(1)}}{\partial \hat{a}^{(1)}} \cdot \frac{\partial \hat{a}^{(1)}}{\partial z^{(1)}} \cdot \frac{\partial z^{(1)}}{\partial w^{(1)}}$$

$$a^{(1)} = \text{Dropout}(\hat{a}^{(1)}, p=0.1) \Rightarrow \frac{\partial a^{(1)}}{\partial \hat{a}^{(1)}} = \begin{bmatrix} .9 & .1 \end{bmatrix} \quad \textcircled{III}$$

Ⓘ , Ⓜ , Ⓢ

$$\frac{\partial L}{\partial w^{(1)}} = \left(\hat{y} w^{(2)} \cdot \begin{bmatrix} .9 & .1 \end{bmatrix} + .1 \cdot \begin{bmatrix} .9 & .1 \end{bmatrix} \right) x^T$$

$$4 \quad y(u, v, z) = \varphi(u, v, z)$$

- Jacobian J is the matrix of first order partial derivatives of y with respect to u, v, z

$$J = \begin{bmatrix} \frac{\partial y}{\partial u} & \frac{\partial y}{\partial v} & \frac{\partial y}{\partial z} \end{bmatrix}$$

- Hessian: is the matrix of second order

$$H = \begin{bmatrix} \frac{\partial^2 y}{\partial u^2} & \frac{\partial^2 y}{\partial u \partial v} & \frac{\partial^2 y}{\partial u \partial z} \\ \frac{\partial^2 y}{\partial v \partial u} & \frac{\partial^2 y}{\partial v^2} & \frac{\partial^2 y}{\partial v \partial z} \\ \frac{\partial^2 y}{\partial z \partial u} & \frac{\partial^2 y}{\partial z \partial v} & \frac{\partial^2 y}{\partial z^2} \end{bmatrix}$$

$$y(u, v, z) \rightarrow \nabla y = \begin{bmatrix} \frac{\partial y}{\partial u} \\ \frac{\partial y}{\partial v} \\ \frac{\partial y}{\partial z} \end{bmatrix} \rightarrow J_{\vec{\nabla} y} = \begin{bmatrix} \frac{\partial \nabla y}{\partial u} & \frac{\partial \nabla y}{\partial v} & \frac{\partial \nabla y}{\partial z} \end{bmatrix}$$

$$J_{\vec{\nabla} y} = \begin{bmatrix} \nabla^T \frac{\partial y}{\partial u} \\ \nabla^T \frac{\partial y}{\partial v} \\ \nabla^T \frac{\partial y}{\partial z} \end{bmatrix}$$

$$J_{\vec{\nabla} y} = H$$

$$J_i = 0.15 \left(y_d - \sum_{k=1}^n \delta_k w_k x_k \right)^2, \quad \delta_k \sim \text{Normal}(1, \sigma^2)$$

$$\frac{\partial J_i}{\partial w_i} = 0.15 \times 2 (-\delta_i x_i) \left(y_d - \sum_{k=1}^n \delta_k w_k x_k \right)$$

$$\frac{\partial J_i}{\partial w_i} = +\delta_i x_i \left(y_d - \sum_{k=1}^n \delta_k w_k x_k \right)$$

Dropout به عنوان Regularization در روش Dropout در هنگام آموزش به صورت تصادفی تعدادی از نودها (دورهای شبکه) را حذف می‌کند. در هر لای drop می‌شوند و از بازی خارج می‌شوند. این به این معناست که در هر مرحله از آموزش برخی از نودها از فرایند آموزش خارج شده و مدل همچنان در شبکه‌ها و پارامترها به یاد و تثبیت می‌کند. روش dropout باعث می‌شود که مدل از یادگیری ویژگی‌های خاص و بیش از حد وابسته به داده‌های آموزش جلوگیری کند. از این جهت overfitting جلوگیری می‌کند و به مدل کمک می‌کند تا به خوبی تعمیم دهد.

Regularized-Non: تابع هدف ترکیبی از دو تابع هدف معمولی و تابع هدف منظم است.

تابع هدف معمولی: این بخش به تعریف کمترین خطای مدل می‌پردازد و این بخش به مدل می‌گوید که به بیشترین بهره‌ای اکتفا کند از داده‌های آموزشی داشته باشد.

تابع هدف منظم: این بخش به منظور جلوگیری از overfitting و بهبود تعمیم‌پذیری مدل به کار می‌رود.

در این بخش معمولاً به صورت (Penalty) به نودهای مدل یا پیچیدگی آن اضافه می‌شود به عنوان مثال از L_2 تابع هدف به صورت زیر:

$$L = \sum_{i=1}^n \text{Loss}(\hat{y}_i, y_i) + \lambda \sum_{j=1}^m w_j^2$$

که ضریب تنظیم است که اندازه جریمه را کنترل می‌کند. w_j وزن نودهای مدل است. این بخش

مدل را مجبور می‌کند برای داده‌های وزن‌های بزرگ که می‌تواند باعث overfitting شود.

حال Regularized-Non: تابع هدف است که ترکیب این بخش به است مبنی بر به وقت

بیشترین بهره‌ای از داده‌های آموزشی به هم که تعمیم‌پذیری مدل به احوال به هم که کنترل می‌کند.

6 $f(x) = g'(x)$ $f(x^*) = 0$, $f'(x^*) \neq 0$ from net

newton's method:

$$f(x_{k+1}) \approx f(x_k) + f'(x_k)t + \frac{1}{2} f''(x_k)t^2$$

$$\frac{df(x_{k+1})}{dt} = 0 \Rightarrow f'(x_k) + f''(x_k)t = 0 \Rightarrow t = -\frac{f'(x_k)}{f''(x_k)}$$

$$x_{k+1} = x_k + t = x_k - \frac{f'(x_k)}{f''(x_k)} \quad \|x_{k+1} - x^*\| \leq \frac{1}{2} \|x_k - x^*\|^2$$

out question: $f(x) = g'(x) \Rightarrow$

$$x^* = \arg \min g(x) \leadsto \frac{\partial g(x)}{\partial x} = 0 \quad f(x^*) = 0$$

$$g(x_k + t) = g(x_k) + g'(x_k)t + \frac{1}{2} g''(x_k)t^2$$

$$\frac{\partial g(x_{k+1})}{\partial t} = 0 \Rightarrow g'(x_k) + \frac{1}{2} \cdot 2t(g''(x_k)) = 0 \Rightarrow t = -\frac{g'(x_k)}{g''(x_k)}$$

$$x_{k+1} = x_k + t \Rightarrow x_{k+1} = x_k - \frac{g'(x_k)}{g''(x_k)}$$

$g'(x) = f(x) \rightarrow$

$x_{k+1} = x_k - \frac{f(x_k)}{f'(x_k)}$

$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$

$$x_{n+1} = x_n - \frac{g'(x_n)}{g''(x_n)}$$

second-order: $e_n = x_n - x^*$, $e_{n+1} = x_{n+1} - x^*$

$$f(x_n) = f(x^* + e_n) = f(x^*) + f'(x^*)e_n + \frac{1}{2} f''(x^*)e_n^2 + \dots$$

$$f(x_n) \approx f'(x^*)e_n + \frac{1}{2} f''(x^*)e_n^2$$

$$f'(x_n) = f'(x^* + e_n) = \underbrace{f'(x^*) + f''(x^*)e_n + \frac{1}{2}f'''(x^*)e_n^2 + \dots}_{f'(x^*) + f''(x^*)e_n + \frac{1}{2}f'''(x^*)e_n^2}$$

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)} \Rightarrow x_{n+1} = x_n - \frac{f(x^*) + f'(x^*)e_n + \frac{1}{2}f''(x^*)e_n^2}{f'(x^*) + f''(x^*)e_n}$$

$$e_{n+1} = x_{n+1} - x^* = \left(x_n - \frac{f(x^*) + f'(x^*)e_n + \frac{1}{2}f''(x^*)e_n^2}{f'(x^*) + f''(x^*)e_n} \right) - x^* \xrightarrow{e_n = x_n - x^*}$$

$$e_{n+1} = e_n - \frac{f'(x^*)e_n + \frac{1}{2}f''(x^*)e_n^2}{f'(x^*) + f''(x^*)e_n} \underset{e_n \text{ small}}{\approx} e_n - \frac{f'(x^*)e_n + \frac{1}{2}f''(x^*)e_n^2}{f'(x^*)}$$

$$e_{n+1} \approx \cancel{e_n} - \cancel{e_n} + \frac{1}{2} \frac{f''(x^*)e_n^2}{f'(x^*)}$$

$$e_{n+1} \approx \underbrace{\frac{1}{2} \cdot \frac{f''(x^*)}{f'(x^*)}}_C e_n^2 \Rightarrow e_{n+1} \approx C |e_n|^2, \quad C = \frac{1}{2} \frac{f''(x^*)}{f'(x^*)}$$

$$\boxed{7} \quad \hat{y}_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}} \quad L(z, y) = - \sum_{k=1}^K y_k \log \hat{y}_k$$

$$\frac{\partial L}{\partial \hat{y}_i} = \frac{\partial}{\partial \hat{y}_i} \left(- \sum_{k=1}^K y_k \log \hat{y}_k \right) = - \frac{y_i}{\hat{y}_i} \quad (\uparrow)$$

$$\frac{\partial \hat{y}_i}{\partial z_j} = \frac{\partial}{\partial z_j} \left(\frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}} \right) = \hat{y}_i (\delta_{ij} - \hat{y}_i)$$

$$\frac{\partial L}{\partial z} = \sum_{j=1}^K \frac{\partial L}{\partial \hat{y}_j} \cdot \frac{\partial \hat{y}_j}{\partial z_i} = \sum_{j=1}^K - \frac{y_j}{\hat{y}_j} \cdot \hat{y}_j (\delta_{ij} - \hat{y}_i)$$

$$\frac{\partial L}{\partial z} = \hat{y}_i - y_i \quad \rightarrow \quad \boxed{\frac{\partial L}{\partial z} = \hat{y} - y}$$

$$i=j \Rightarrow \frac{\partial^2 L}{\partial z_i^2} = \hat{y}_i (1 - \hat{y}_i)$$

(ب)

$$i \neq j \Rightarrow \frac{\partial^2 L}{\partial z_i \partial z_j} = -\hat{y}_i \hat{y}_j \quad \frac{\partial L}{\partial z_i}$$

(الف)

$$H_{ij} = \frac{\partial^2 L}{\partial z_i \partial z_j} = \frac{\partial}{\partial z_j} (\hat{y}_i - y_i)$$

$$H_{ij} = \begin{cases} \hat{y}_i (1 - \hat{y}_i) & i=j \\ -\hat{y}_i \hat{y}_j & i \neq j \end{cases}$$

$$H = \begin{bmatrix} \hat{y}_1 (1 - \hat{y}_1) & -\hat{y}_1 \hat{y}_2 & \dots & -\hat{y}_1 \hat{y}_K \\ -\hat{y}_2 \hat{y}_1 & \hat{y}_2 (1 - \hat{y}_2) & \dots & -\hat{y}_2 \hat{y}_K \\ \vdots & \vdots & \ddots & \vdots \\ -\hat{y}_K \hat{y}_1 & -\hat{y}_K \hat{y}_2 & \dots & \hat{y}_K (1 - \hat{y}_K) \end{bmatrix}$$

(ب) باید ثابت شود $v^T H v > 0$

$$H_{ij} = \begin{cases} y_i (1 - y_i) & i = j \\ -\hat{y}_i \hat{y}_j & i \neq j \end{cases}$$

non-negative

$$v^T H v = \sum_{i=1}^K \sum_{j=1}^K v_i v_j H_{ij}$$

covariances between classes

$$v^T H v = \sum_{i=1}^K v_i^2 \hat{y}_i (1 - \hat{y}_i) - \sum_{i=1}^K \sum_{j \neq i} v_i v_j \hat{y}_i \hat{y}_j$$

sum of variance $\rightarrow > 0$

$$v^T H v > 0$$

خاصیتی: واریانس افعال هر کلاس، مثبت می‌دهد.

خاصیتی: کواریانس بین کلاس‌های مختلف، منفی می‌دهد.

☆ مثبت یعنی بودن این ماتریس قطری مثبت است و این ماتریس کواریانس دارد زیرا عناصر احتمالات تابع softmax منفی می‌دهد. مجموع آن 1 را هم می‌بینیم یعنی 1 مقدار دارند. ماتریس کواریانس همواره مثبت و منفی می‌دهد. این معنای هم مقدار و اثر آنها نیز منفی است.

ج) مثبت یعنی بودن ماتریس هین نشان می‌دهد که تابع لگاریتمی متقاطع نسبت به ورودی‌های فعال از softmax محسوب است. محاسبه بودن این تابع تضمین می‌کند که به رفتار شبکه یادگیری روی این تابع به سمت کمینه از تابع می‌کند و این ویژگی به بهبود یادگیری و دقت در یادگیری شبکه کمک می‌کند.

بخش عملی

Github link:

[Google drive link:](#)

<https://drive.google.com/drive/folders/1LhpeUO8tcGSwovNlqIx1e3SonjVImZYa?usp=sharing>

سوال ۱:

برای پیاده سازی rosenbrock کد زیر را اجرا میکنیم که در آن مشتق اول نسبت به $x[0]$, $x[1]$ میگیریم

```
def rosenbrock(x):
    """Returns the value and gradient of Rosenbrock's function at x: 2d vector"""
    x1, x2 = x[0], x[1]
    val = 100*(x[1]-x[0]**2)**2+(x[0]-1)**2
    dv_dx0 = -400 * x[0] * (x[1] - x[0]**2) + 2 * (x[0] - 1)
    dv_dx1 = 200 * (x[1] - x[0]**2)
    grad = np.array([dv_dx0, dv_dx1])

    return val, grad
```

برای پیاده سازی rosenbrock_hessian کد زیر را اجرا میکنیم که در آن مشتق دوم نسبت به $x[0]$, $x[1]$ میگیریم

```
def rosenbrock_hessian(x):
    """Returns the value, gradient and hessian of Rosenbrock's function at x: 2d vector"""
    val, grad = rosenbrock(x)

    # Hessian matrix components
    d2v_dx0x0 = 1200 * x[0]**2 - 400 * x[1] + 2
    d2v_dx0x1 = -400 * x[0]
    d2v_dx1x0 = -400 * x[0]
    d2v_dx1x1 = 200

    # Hessian matrix
    hessian = np.array([[d2v_dx0x0, d2v_dx0x1],
                        [d2v_dx1x0, d2v_dx1x1]])

    return val, grad, hessian
```

و GD هم به صورت زیر پیاده سازی می شد

```
def GD(f, theta0, alpha, stop_tolerance=1e-10, max_steps=1000000):
    """Runs gradient descent algorithm on f.

    Args:
        f: function that when evaluated on a Theta of same dtype and shape as Theta0
            returns a tuple (value, dv_dtheta) with dValuedTheta of the same shape
            as Theta
        theta0: starting point
        alpha: step length
        stop_tolerance: stop iterations when improvement is below this threshold
        max_steps: maximum number of steps

    Returns:
```

```

    tuple:
    - theta: optimum theta found by the algorithm
    - history: list of length num_steps containing tuples (theta, (val,
dv_dtheta: np.array))

"""
history = []

theta = theta0

step = 0

    # Initial function evaluation
val, grad = f(theta)
print(grad)
history.append((theta.copy(), (val, grad)))

while step < max_steps:
    # Gradient descent step: update theta
    theta = theta - alpha * grad

    # Evaluate function at new theta
    new_val, grad = f(theta)
    history.append((theta.copy(), (new_val, grad)))

    # Check for convergence
    improvement = abs(val - new_val)
    if improvement < stop_tolerance:
        break

    # Update the value for next iteration
    val = new_val
    step += 1

history.append([theta, f(theta)])
return theta, history

```

برای پیدا کردن optimum تابع rosenbrock داریم:

```

X0 = [0.,2.]
Xopt, Xhist = GD(rosenbrock, X0, alpha=1e-3, stop_tolerance=1e-10, max_steps=1e6)

print ("Found optimum at %s in %d steps (true minimum is at [1,1])" % (Xopt,
len(Xhist)))

# Plot how the value changes over iterations
#TODO
# Extract function values over iterations
values = [entry[1][0] for entry in Xhist] # Get function values from history


# Plot the convergence of function value

```

```
plt.figure(figsize=(10, 6))
plt.plot(values, label="Rosenbrock function value")
plt.yscale("log") # Use log scale to see the improvement more clearly
plt.xlabel("Iterations")
plt.ylabel("Function Value (log scale)")
plt.title("Convergence of Gradient Descent on Rosenbrock Function")
plt.legend()
plt.grid(True)
plt.show()
```

مقدار اپتیموم [0.99964674 0.99929219] است که به مقدار ۱ و ۱ نزدیک می باشد و نمودار آن هم به شکل زیر است.

Found optimum at [0.99964674 0.99929219] in 16855 steps (true minimum is at [1,1])



برای پیاده سازی متد Newton به اپتیموم [1,1] میشود

Found optimum at [1. 1.] (true minimum is at [1,1])

```
# Newton's Method
def Newton(f, theta0, alpha=1, stop_tolerance=1e-10, max_steps=1000000):
    """Performs Newton's optimization method with a simple line search.

    Args:
        f: function that when evaluated on a Theta of same dtype and shape as Theta0
            returns a tuple (value, gradient, hessian), where gradient and Hessian
            have the same shape as Theta.
        theta0: starting point.
        alpha: step length for backtracking line search (default is 1).
        stop_tolerance: stop iterations when the norm of the gradient is below this
            threshold.
        max_steps: maximum number of iterations.
```

```

Returns:
    tuple:
        - theta: optimal Theta after convergence or maximum steps.
        - history: list of tuples (theta, value, gradient) containing the
optimization path.
"""
theta = theta0
history = []
# TODO
for step in range(max_steps):
    # Evaluate the function, gradient, and Hessian
    val, grad, hessian = f(theta)

    # Store the current theta and function value
    history.append((theta.copy(), val))

    # Check if gradient norm is below the tolerance
    if np.linalg.norm(grad) < stop_tolerance:
        break

    # Update rule using inverse of the Hessian
    try:
        theta -= np.linalg.inv(hessian).dot(grad)
    except np.linalg.LinAlgError:
        print("Hessian is singular at step", step)
        break # Stop if the Hessian is singular

return theta, history

# Test Newton's method on the Rosenbrock function
X0 = [0., 2.] # Initial guess
Xopt, Xhist = Newton(rosenbrock_hessian, X0)

print("Found optimum at %s (true minimum is at [1,1])" % Xopt)

```

Part two: MLP for MNIST Classification

برای این بخش ۳ فایل مد نظر را کامل کرده و داریم: با بچ سایز ۳۲ و تعداد epoch=20
نتایج فاز ترین:

و با مدل mlp با لایه های FCLayer and SigmoidLayer:

```

sigmoidMLP = nn.Sequential(
    FCLayer(784, 128),
    SigmoidLayer(),
    FCLayer(128, 10)
)

```

Epoch [1] Average training loss: 0.0771, Average training accuracy: 0.4956

Epoch [2] Average training loss: 0.0556, Average training accuracy: 0.7378

Epoch [3] Average training loss: 0.0506, Average training accuracy: 0.7836

Epoch [4] Average training loss: 0.0482, Average training accuracy: 0.8046

Epoch [5] Average training loss: 0.0466, Average training accuracy: 0.8173
Epoch [6] Average training loss: 0.0456, Average training accuracy: 0.8254
Epoch [7] Average training loss: 0.0448, Average training accuracy: 0.8307
Epoch [8] Average training loss: 0.0442, Average training accuracy: 0.8344
Epoch [9] Average training loss: 0.0437, Average training accuracy: 0.8384
Epoch [10] Average training loss: 0.0432, Average training accuracy: 0.8408
Epoch [11] Average training loss: 0.0429, Average training accuracy: 0.8430
Epoch [12] Average training loss: 0.0426, Average training accuracy: 0.8455
Epoch [13] Average training loss: 0.0423, Average training accuracy: 0.8465
Epoch [14] Average training loss: 0.0420, Average training accuracy: 0.8474
Epoch [15] Average training loss: 0.0418, Average training accuracy: 0.8489
Epoch [16] Average training loss: 0.0415, Average training accuracy: 0.8504
Epoch [17] Average training loss: 0.0413, Average training accuracy: 0.8517
Epoch [18] Average training loss: 0.0411, Average training accuracy: 0.8514
Epoch [19] Average training loss: 0.0409, Average training accuracy: 0.8529
Epoch [20] Average training loss: 0.0407, Average training accuracy: 0.8533

برای فاز تست داریم :

The test accuracy is 0.8628.

برای MLP با لایه های FCLayer and ReLULayer:

```
reluMLP = nn.Sequential(  
    FCLayer(784, 128),  
    ReLULayer(),  
    FCLayer(128, 10)  
)
```

Epoch [1] Average training loss: 0.0712, Average training accuracy: 0.6373
Epoch [2] Average training loss: 0.0467, Average training accuracy: 0.8231
Epoch [3] Average training loss: 0.0399, Average training accuracy: 0.8589
Epoch [4] Average training loss: 0.0357, Average training accuracy: 0.8765
Epoch [5] Average training loss: 0.0327, Average training accuracy: 0.8874
Epoch [6] Average training loss: 0.0305, Average training accuracy: 0.8953
Epoch [7] Average training loss: 0.0288, Average training accuracy: 0.9011
Epoch [8] Average training loss: 0.0274, Average training accuracy: 0.9060

Epoch [9] Average training loss: 0.0262, Average training accuracy: 0.9103
 Epoch [10] Average training loss: 0.0253, Average training accuracy: 0.9131
 Epoch [11] Average training loss: 0.0244, Average training accuracy: 0.9162
 Epoch [12] Average training loss: 0.0237, Average training accuracy: 0.9185
 Epoch [13] Average training loss: 0.0231, Average training accuracy: 0.9205
 Epoch [14] Average training loss: 0.0225, Average training accuracy: 0.9227
 Epoch [15] Average training loss: 0.0220, Average training accuracy: 0.9248
 Epoch [16] Average training loss: 0.0215, Average training accuracy: 0.9264
 Epoch [17] Average training loss: 0.0211, Average training accuracy: 0.9279
 Epoch [18] Average training loss: 0.0207, Average training accuracy: 0.9289
 Epoch [19] Average training loss: 0.0204, Average training accuracy: 0.9297
 Epoch [20] Average training loss: 0.0200, Average training accuracy: 0.9309

برای فاز تست داریم :

The test accuracy is 0.9312.

در این تغییرات، مدل شبکه عصبی را طوری تغییر دادم که احتمال اورفیت (Overfitting) بیشتری داشته باشد. اورفیت زمانی اتفاق می افتد که مدل بیش از حد پیچیده می شود و بیشتر به حفظ جزئیات داده های آموزشی پرداخته و قادر به تعمیم خوب روی داده های جدید نخواهد بود.

تغییرات اصلی که اعمال کردم:

- افزایش پیچیدگی شبکه (لایه های بیشتر و تعداد نوروں های بیشتر):
- من تعداد نوروں ها در لایه های مخفی را افزایش دادم (مثلاً از ۵۱۲ به ۱۰۲۴ و از ۲۵۶ به ۵۱۲) که باعث پیچیده تر شدن مدل می شود.
- همچنین تعداد لایه ها را از ۳ لایه به ۵ لایه افزایش دادم. این باعث می شود که مدل ظرفیت بیشتری برای یادگیری پارامترها داشته باشد.

۲. حذف لایه های Dropout:

- در مدل اصلی، از لایه های Dropout استفاده می شود که یک روش منظم کننده است و به مدل کمک می کند از اورفیت جلوگیری کند.
- من این لایه های Dropout را حذف کردم تا مدل بدون هیچ نوع محدودیتی بتواند پارامترها را یاد بگیرد و بیشتر به حفظ جزئیات داده های آموزشی بپردازد. این تغییر باعث می شود مدل به راحتی روی داده های آموزشی اورفیت کند.

۳. افزایش ظرفیت مدل:

- با افزایش تعداد لایه ها و نوروں ها، مدل به طور کلی پیچیده تر و بزرگتر شده است. این امر باعث می شود که مدل توانایی یادگیری بیشتر و به تبع آن احتمال اورفیت شدن نیز افزایش یابد، به ویژه زمانی که داده ها کافی نباشند یا مدل برای تعداد زیادی از دوره های آموزشی آموزش ببیند.

چرا این تغییرات باعث اورفیت می شوند؟

- **افزایش تعداد پارامترها:** مدل حالا تعداد بیشتری پارامتر برای یادگیری دارد. این باعث می شود مدل قادر به یادگیری جزئیات بیشتری از داده های آموزشی باشد و به راحتی روی داده های آموزشی اورفیت کند.
- **حذف Dropout:** Dropout یک تکنیک است که برای جلوگیری از اورفیت استفاده می شود. حذف این تکنیک باعث می شود که مدل به صورت کامل به یادگیری داده ها پرداخته و احتمال اورفیت شدن بیشتر می شود.
- **شبکه عمیق تر:** با اضافه کردن لایه های بیشتر، مدل پیچیده تر شده و ظرفیت یادگیری آن افزایش می یابد که می تواند باعث اورفیت روی داده های آموزشی شود.

نتیجه:

این تغییرات باعث می‌شود که مدل توانایی یادگیری جزئیات زیادی از داده‌های آموزشی داشته باشد و احتمالاً نتواند به خوبی روی داده‌های جدید عمل کند (یعنی اورفیت می‌کند). این مدل برای آزمایش اورفیت مناسب است، به خصوص اگر داده‌های آموزشی کوچک یا تعداد دوره‌های آموزش زیاد باشد.

لایه با اورفیت :

```
#TODO: overfit the reluMLP model
num_epoch = 50

reluMLP = nn.Sequential(
    FCLayer(784, 1024),
    ReLULayer(),
    FCLayer(1024, 512),
    ReLULayer(),
    FCLayer(512, 256),
    ReLULayer(),
    FCLayer(256, 128),
    ReLULayer(),
    FCLayer(128, 64),
    ReLULayer(),
    FCLayer(64, 10)
)
criterion = nn.MSELoss()

# Initialize optimizer
sgd = SGD(reluMLP.parameters(), learning_rate=0.5)

# Train the model
reluMLP = train(reluMLP, criterion, sgd, train_dataloader, num_epoch, device=device)

test(reluMLP, test_dataloader, device)
```

```
Epoch [1] Average training loss: 0.0154, Average training accuracy: 0.9286
Epoch [2] Average training loss: 0.0061, Average training accuracy: 0.9726
Epoch [3] Average training loss: 0.0041, Average training accuracy: 0.9831
Epoch [4] Average training loss: 0.0029, Average training accuracy: 0.9884
Epoch [5] Average training loss: 0.0022, Average training accuracy: 0.9921
Epoch [6] Average training loss: 0.0016, Average training accuracy: 0.9946
Epoch [7] Average training loss: 0.0012, Average training accuracy: 0.9966
Epoch [8] Average training loss: 0.0009, Average training accuracy: 0.9974
Epoch [9] Average training loss: 0.0007, Average training accuracy: 0.9981
Epoch [10] Average training loss: 0.0006, Average training accuracy: 0.9989
```

[illegible]

Epoch [43] Average training loss: 0.0000, Average training accuracy: 1.0000
Epoch [44] Average training loss: 0.0000, Average training accuracy: 1.0000
Epoch [45] Average training loss: 0.0000, Average training accuracy: 1.0000
Epoch [46] Average training loss: 0.0000, Average training accuracy: 1.0000
Epoch [47] Average training loss: 0.0000, Average training accuracy: 1.0000
Epoch [48] Average training loss: 0.0000, Average training accuracy: 1.0000
Epoch [49] Average training loss: 0.0000, Average training accuracy: 1.0000
Epoch [50] Average training loss: 0.0000, Average training accuracy: 1.0000

از epoch حدوداً ۱۰ به بعد داده های آموزش را حفظ کرده است

The test accuracy is 0.9836.

بعد از اضافه کردن دراپ اوت:

```
from layers import DropoutLayer
#TODO: add DropoutLayer to your model

#TODO: overfit the reluMLP model
num_epoch = 30

reluMLP = nn.Sequential(
    FCLayer(784, 1024),
    ReLULayer(),
    DropoutLayer(0.5), # Dropout layer with a rate of 0.5

    FCLayer(1024, 512),
    ReLULayer(),
    DropoutLayer(0.5), # Dropout layer with a rate of 0.5

    FCLayer(512, 256),
    ReLULayer(),
    DropoutLayer(0.5), # Dropout layer with a rate of 0.5
    FCLayer(256, 128),
    ReLULayer(),
    DropoutLayer(0.5), # Dropout layer with a rate of 0.5

    FCLayer(128, 64),
    ReLULayer(),
    DropoutLayer(0.5), # Dropout layer with a rate of 0.5

    FCLayer(64, 10)
)
criterion = nn.MSELoss()

# Initialize optimizer
```



```
sgd = SGD(reluMLP.parameters(), learning_rate=0.5)

# Train the model
reluMLP = train(reluMLP, criterion, sgd, train_dataloader, num_epoch, device=device)

test(reluMLP, test_dataloader, device)
```

```
Epoch [1] Average training loss: 0.0875, Average training accuracy: 0.1939
Epoch [2] Average training loss: 0.0793, Average training accuracy: 0.2974
Epoch [3] Average training loss: 0.0737, Average training accuracy: 0.3886
Epoch [4] Average training loss: 0.0673, Average training accuracy: 0.4494
Epoch [5] Average training loss: 0.0636, Average training accuracy: 0.4680
Epoch [6] Average training loss: 0.0613, Average training accuracy: 0.4785
Epoch [7] Average training loss: 0.0601, Average training accuracy: 0.4808
Epoch [8] Average training loss: 0.0592, Average training accuracy: 0.4874
Epoch [9] Average training loss: 0.0589, Average training accuracy: 0.4880
Epoch [10] Average training loss: 0.0585, Average training accuracy: 0.4889
Epoch [11] Average training loss: 0.0582, Average training accuracy: 0.4918
Epoch [12] Average training loss: 0.0580, Average training accuracy: 0.4896
Epoch [13] Average training loss: 0.0580, Average training accuracy: 0.4914
Epoch [14] Average training loss: 0.0576, Average training accuracy: 0.4942
Epoch [15] Average training loss: 0.0577, Average training accuracy: 0.4949
Epoch [16] Average training loss: 0.0576, Average training accuracy: 0.4940
Epoch [17] Average training loss: 0.0575, Average training accuracy: 0.4929
Epoch [18] Average training loss: 0.0573, Average training accuracy: 0.4942
Epoch [19] Average training loss: 0.0572, Average training accuracy: 0.4963
Epoch [20] Average training loss: 0.0572, Average training accuracy: 0.4947
Epoch [21] Average training loss: 0.0571, Average training accuracy: 0.4967
Epoch [22] Average training loss: 0.0570, Average training accuracy: 0.4949
Epoch [23] Average training loss: 0.0568, Average training accuracy: 0.4970
Epoch [24] Average training loss: 0.0567, Average training accuracy: 0.4974
Epoch [25] Average training loss: 0.0568, Average training accuracy: 0.4980
Epoch [26] Average training loss: 0.0566, Average training accuracy: 0.4986
Epoch [27] Average training loss: 0.0566, Average training accuracy: 0.4994
```

Epoch [28] Average training loss: 0.0566, Average training accuracy: 0.4994

Epoch [29] Average training loss: 0.0566, Average training accuracy: 0.4999

Epoch [30] Average training loss: 0.0564, Average training accuracy: 0.4998

The test accuracy is 0.5041.

Introduction to Loss Functions

```

class SimpleMLP(nn.Module):
    def __init__(self, input_dim, hidden_dim, output_dim, num_hidden_layers=1,
last_layer_activation_fn=None):
        super(SimpleMLP, self).__init__()
        # TODO: Define the layers of the MLP
        layers = []

        # Input layer
        layers.append(nn.Linear(input_dim, hidden_dim))
        layers.append(nn.ReLU())

        # Hidden layers
        for _ in range(num_hidden_layers - 1):
            layers.append(nn.Linear(hidden_dim, hidden_dim))
            layers.append(nn.ReLU())

        # Output layer
        layers.append(nn.Linear(hidden_dim, output_dim))

        # Add the last layer activation function only if it's provided
        if last_layer_activation_fn is not None:
            layers.append(last_layer_activation_fn())

        # Combine layers into a Sequential module
        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)

class SimpleMLPTrainer:
    def __init__(self, model, criterion, optimizer):
        self.model = model
        self.criterion = criterion
        self.optimizer = optimizer

    def train(self, train_loader, num_epochs):
        #TODO: Implement the training loop
        #Note: You should also print the training loss at each epoch, use tqdm for
progress bar
        #Note: You should return the training loss at each epoch
        self.model.train()
        epoch_losses = []

        for epoch in range(num_epochs):
            total_loss = 0.0
            for inputs, targets in tqdm(train_loader, desc=f"Epoch
{epoch+1}/{num_epochs}"):
                # Ensure targets are 1D (class indices)

```

```

        targets = targets.view(-1)  # Convert to 1D tensor if needed

        # Forward pass with log_softmax for NLLLoss
        outputs = self.model(inputs)
        loss = self.criterion(outputs, targets)

        self.optimizer.zero_grad()
        loss.backward()
        self.optimizer.step()

        total_loss += loss.item()

    average_loss = total_loss / len(train_loader)
    epoch_losses.append(average_loss)
    print(f"Epoch {epoch+1}/{num_epochs}, Loss: {average_loss:.4f}")

    return epoch_losses
pass

def evaluate(self, val_loader):
    #TODO: Implement the evaluation loop
    #Note: You should return the validation loss and accuracy
    self.model.eval()
    total_loss = 0.0
    correct = 0
    total = 0
    with torch.no_grad():
        for inputs, targets in val_loader:
            # Ensure targets are 1D (class indices)
            targets = targets.view(-1)

            # Forward pass with log_softmax for NLLLoss
            outputs = F.log_softmax(self.model(inputs), dim=1)
            loss = self.criterion(outputs, targets)
            total_loss += loss.item()

            # Calculate accuracy
            _, predicted_classes = torch.max(outputs, 1)
            correct += (predicted_classes == targets).sum().item()
            total += targets.size(0)

    average_loss = total_loss / len(val_loader)
    accuracy = (correct / total) * 100 if total > 0 else 0
    return average_loss, accuracy

pass

```

```

# Load dataset
train_url =
"https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv"
data = pd.read_csv(train_url)

# Preprocessing (simple example)
data = data[['Pclass', 'Sex', 'Age', 'Fare', 'Survived']].dropna()
data['Sex'] = data['Sex'].map({'male': 0, 'female': 1})

# TODO: Convert the data to PyTorch tensors and create a DataLoader
X = data[['Pclass', 'Sex', 'Age', 'Fare']].values
y = data['Survived'].values

scaler = StandardScaler()
X = scaler.fit_transform(X)
X_tensor = torch.tensor(X, dtype=torch.float32)
y_tensor = torch.tensor(y, dtype=torch.float32).view(-1, 1) # Reshape for
compatibility

# TODO: Split the data into training and validation sets
dataset = TensorDataset(X_tensor, y_tensor)
train_size = int(0.8 * len(dataset))
val_size = len(dataset) - train_size
train_dataset, val_dataset = random_split(dataset, [train_size, val_size])
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=32)

# TODO: Define the model, criterion, and optimizer
input_dim = X.shape[1] # Number of features
hidden_dim = 16 # Adjust based on experimentation
output_dim = 1 # Binary classification (Survived or not)

# Model
model = SimpleMLP(input_dim, hidden_dim, output_dim, num_hidden_layers=2,
last_layer_activation_fn=None)

# Criterion (loss function) and optimizer
criterion = nn.BCEWithLogitsLoss() # Use BCEWithLogitsLoss for binary classification
without sigmoid activation in the model
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Print dataset and model information
print(f"Training samples: {len(train_dataset)}, Validation samples:
{len(val_dataset)}")
print(model)

```

```

Training samples: 571, Validation samples: 143
SimpleMLP(
  (model): Sequential(
    (0): Linear(in_features=4, out_features=16, bias=True)
    (1): ReLU()
  )
)

```



```

        (2): Linear(in_features=16, out_features=16, bias=True)
        (3): ReLU()
        (4): Linear(in_features=16, out_features=1, bias=True)
    )
)

```

L1Loss

```

from torch.nn import L1Loss

# TODO: Train the model

model = SimpleMLP(input_dim=X.shape[1], hidden_dim=16, output_dim=1)
criterion = nn.L1Loss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

trainer = SimpleMLPTrainer(model, criterion, optimizer)
train_losses = trainer.train(train_loader, num_epochs=20)

# TODO: Evaluate the model
validation_loss, validation_accuracy = trainer.evaluate(val_loader)
print(f'Validation Loss: {validation_loss:.4f}, Accuracy:
{validation_accuracy:.2f}%')
Validation Loss: 0.4058, Accuracy: 60.14%

```

MSELoss

```

from torch.nn import MSELoss

# TODO: Train the model
criterion = nn.MSELoss()
optimizer = Adam(model.parameters(), lr=0.01)
trainer = SimpleMLPTrainer(model, criterion, optimizer)
train_losses = trainer.train(train_loader, num_epochs=20)

# TODO: Evaluate the model

print("\nEvaluating the model on the validation set:")
validation_loss, validation_accuracy = trainer.evaluate(val_loader)

print(f"\nValidation Loss: {validation_loss:.4f}")
print(f"Validation Accuracy: { validation_accuracy:.2f}%")
Validation Loss: 0.4058
Validation Accuracy: 60.14%

```

NLLLoss

```

# Run with relu activation function
from torch.nn import NLLLoss

criterion = nn.NLLLoss()
optimizer = Adam(model.parameters(), lr=0.01)
trainer = SimpleMLPTrainer(model, criterion, optimizer)

```

```
# Train the model
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True) # Replace with
your dataset
train_losses = trainer.train(train_loader, num_epochs=20)

# Evaluate the model
print("\nEvaluating the model on the validation set:")
validation_loss, validation_accuracy = trainer.evaluate(val_loader)

print(f"\nValidation Loss: {validation_loss:.4f}")
print(f"Validation Accuracy: {validation_accuracy:.2f}%")
```

بخش Regularization in Machine Learning

```
from sklearn.neural_network import MLPClassifier
```

MLPClassifier را فراخوانی میکنیم
دیتا را لود کرده و به داده تست و ترین جدا میکنیم:

```
# 1. Load and Prepare the Iris Dataset
iris = load_iris()
X = iris.data # Features
y = iris.target # Target labels

# Select only two classes for binary classification (Setosa and Versicolor)
binary_mask = y < 2
X, y = X[binary_mask], y[binary_mask]
# Select two features for 2D visualization (Sepal Length and Petal Length)
X = X[:, [0, 2]]
# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42)
# Standardize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

توابع `plot_decision_boundary` و `create_decision_boundary_gif` را کامل کرده و gif را
در پوشه ذخیره کردم

```
def plot_decision_boundary(model, X, y, alpha):
    # Define the grid (use meshgrid)
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                          np.arange(y_min, y_max, 0.01))

    # Predict over the grid
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    # Create a figure
```

```

fig, ax = plt.subplots(figsize=(6, 5))

# Plot the decision boundary
ax.contourf(xx, yy, Z, alpha=0.3, levels=[-0.1, 0.1, 1.1], colors=['blue',
'red'])

# Scatter plot of the training data
scatter = ax.scatter(
    X[:, 0], X[:, 1], c=y, cmap='bwr', edgecolor='k', s=50
)

# Title and labels
ax.set_title(f'MLP Decision Boundary (alpha={alpha})')
ax.set_xlabel('Sepal Length (standardized)')
ax.set_ylabel('Petal Length (standardized)')

# Remove axes for clarity
ax.set_xticks([])
ax.set_yticks([])

# Tight layout
plt.tight_layout()

# Save the plot to a BytesIO object
buf = BytesIO()
plt.savefig(buf, format='png')
plt.close(fig)
buf.seek(0)
return Image.open(buf)
def create_decision_boundary_gif(alpha_values, X_train, y_train, n_neurons):

# List to store images
images = []

for idx, alpha in enumerate(alpha_values):
    print(f"Processing alpha={alpha:.4f} ({idx + 1}/{len(alpha_values)})")

    # Create and train the MLP
    mlp = MLPClassifier(hidden_layer_sizes=(n_neurons,), alpha=alpha,
max_iter=1000, random_state=42)
    mlp.fit(X_train, y_train)

    # Plot decision boundary and get the image
    img = plot_decision_boundary(mlp, X_train, y_train, alpha)
    images.append(img)

# Save the images as a GIF
gif_filename = 'mlp_classification_boundaries.gif'
images[0].save(
    gif_filename,

```

```

        save_all=True,
        append_images=images[1:],
        duration=500,
        loop=0
    )

    print(f"GIF saved as '{gif_filename}'")

    # return the gif
    return gif_filename

```

```

# Use np.logspace to generate alpha values, with at least 20 values
alpha_values = np.logspace(-3, 3, 20) # Range from 0.001 to 1000, with 20 steps
# Define the number of neurons in the hidden layer
n_neurons = 10 # This can be adjusted based on the desired model complexity

# Create the decision boundary GIF
gif_dir = create_decision_boundary_gif(alpha_values, X_train, y_train, n_neurons)

```

```

Processing alpha=0.0010 (1/20)
Processing alpha=0.0021 (2/20)
Processing alpha=0.0043 (3/20)
Processing alpha=0.0089 (4/20)
Processing alpha=0.0183 (5/20)
Processing alpha=0.0379 (6/20)
Processing alpha=0.0785 (7/20)
Processing alpha=0.1624 (8/20)
Processing alpha=0.3360 (9/20)
Processing alpha=0.6952 (10/20)
Processing alpha=1.4384 (11/20)
Processing alpha=2.9764 (12/20)
Processing alpha=6.1585 (13/20)
Processing alpha=12.7427 (14/20)
Processing alpha=26.3665 (15/20)
Processing alpha=54.5559 (16/20)
Processing alpha=112.8838 (17/20)
Processing alpha=233.5721 (18/20)
Processing alpha=483.2930 (19/20)
Processing alpha=1000.0000 (20/20)
GIF saved as 'mlp_classification_boundaries.gif'

```